

Simple example of parallel code and how to run it on a cluster

Bruno Juliá-Díaz

U. Barcelona

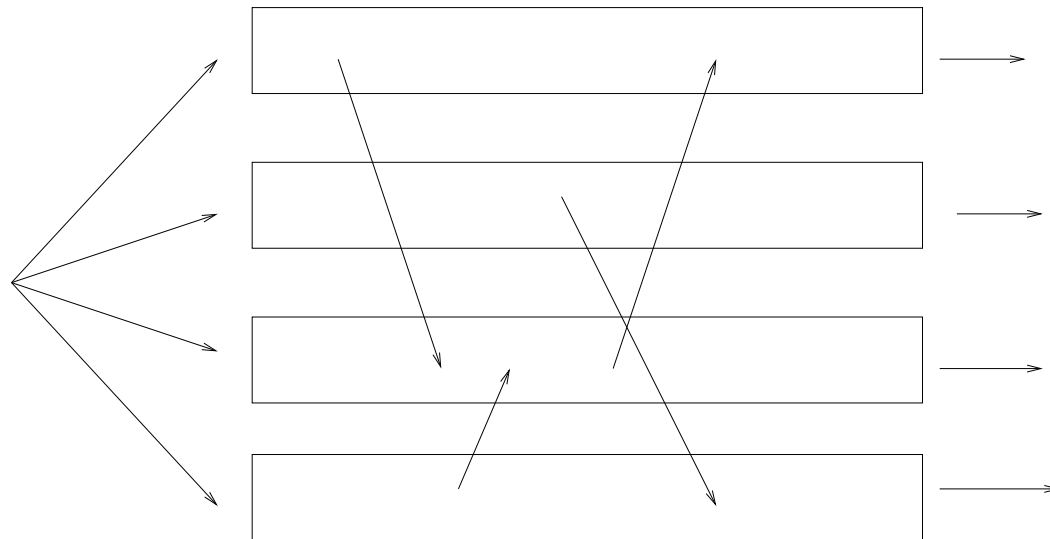
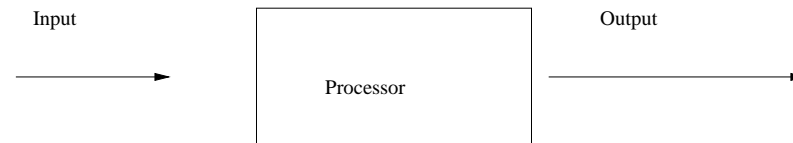
Motivation

- Running in parallel
- one (very) simple example
- a second example with some physics

Running in parallel

Usually we run using one single processor

But we can also run in **several of them** at the same time!



Running in parallel (II)

- Tricky points:
 - The **SAME** code will run in ALL processors
 - How do we make that each of them computes different things?
 - in other words, **how do we tell the code in which processor he is?**
 - Maybe some information should be shared between processors
 - Example: **They need to share preliminar results from time to time**
- We need MPI (*message passing interface*)

Running in parallel (III)

- I'll use MPI with fortran (or *c* should be ok)
- This is quite standard
- you can run in the Barcelona Supercomputer (BSC/CNS)
- on a linux cluster
- on your dual processor (using <http://open-mpi.org>)

Very Simple Example (I)

Very simple parallelization \Rightarrow Simple loop structures:

```
s.f          Mon Dec 18 15:32:08 2006          1

do 2 i=1,2
  E=real(i)
  call computecrosssection(E,CST)
  print*,CST
2 continue
end
subroutine computecrosssection(E,CST)
  CST=E**3
end
```

We can easily parallelize it as each E value is independent from the other E s

Very Simple example (II)

sp.f

Mon Dec 18 15:27:22 2006

1

```
include "mpif.h"
call MPI_INIT(ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD, myid, ierr)
call MPI_COMM_SIZE(MPI_COMM_WORLD, numprocs, ierr)
```

```
do 2 i=1,2
  if (i.ne.myid+1) goto 2
  E=real(i)
  call computecrosssection(E,CST)
  print*,"I am node", myid, "my value",CST
2  continue
CALL MPI_FINALIZE(IERR)
end
```

```
subroutine computecrosssection(E,CST)
  CST=E**3
end
```

Very Simple example (III)

The commands we have used are:

- `include "mpif.h"` Loads MPI library
- `call MPI_COMM_INIT(ierr)` Initialize MPI
- `call MPI_COMM_RANK(MPI_COMM_WORLD, myid, ierr)`
gives the processor number, from 0 to Numprocs-1
- `call MPI_COMM_SIZE(MPI_COMM_WORLD, numprocs, ierr)`
gives the number of processors
- `call MPI_FINALIZE(ierr)`
Finalizes

Very Simple example (IV)

- **key point** We used a conditional to compute ONLY in a selected processor:
if (i.ne.myid+1) goto 2

Four useful commands

Basic sending/receiving commands:

- CALL MPI_SEND(buff, count, MPI_TYPE, dest, tag, comm, ierr)
 - CALL MPI_SEND(ztsend,1,MPI_DOUBLE_COMPLEX,0,n1tag, MPI_COMM_WORLD,MPIERR)
- call MPI_RECV(rbuf, count, MPI_TYPE, source, tag, comm, status, ierr)
 - call MPI_RECV(ztrec,1,MPI_DOUBLE_COMPLEX,IRANK,n1tag, MPI_COMM_WORLD,MPISTAT,MPIERR)
- call MPI_BCAST(buff, count, MPI_TYPE, root, comm, ierr)
 - call MPI_BCAST(ct,1,MPI_DOUBLE_PRECISION,0,MPI_COMM_WORLD,ierr)
- CALL MPI_REDUCE(sendbuf, recbuf, count, MPI_TYPE, MPI_OP, root, comm, ierr)
 - call MPI_REDUCE(xmychi2,ct,1,MPI_DOUBLE_PRECISION,MPI_SUM,0, MPI_COMM_WORLD,ierr)

Executing the code

pa Mon Dec 18 15:28:06 2006 1

```
bruno@melon:~/textos/barna_parallel06> more sp.f
```

```
include "mpif.h"
call MPI_INIT(ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD, myid, ierr)
call MPI_COMM_SIZE(MPI_COMM_WORLD, numprocs, ierr)
```

```
do 2 i=1,2
if (i.ne.myid+1) goto 2
E=real(i)
call compute_crossection(E,CST)
print*, "I am node", myid, "my value", CST
2 continue
CALL MPI_FINALIZE(IERR)
end
```

```
subroutine compute_crossection(E,CST)
CST=E**3
end
```

```
bruno@melon:~/textos/barna_parallel06> mpif77 sp.f
```

```
ifc: warning: The Intel Fortran driver is now named ifort. You can suppress this message
with '-quiet'
```

```
bruno@melon:~/textos/barna_parallel06> mpirun -np 2 a.out
```

```
I am node            1 my value    8.000000
I am node            0 my value    1.000000
```

```
bruno@melon:~/textos/barna_parallel06>
```

Self promotion

I am interested together with a group at Jefferson Lab

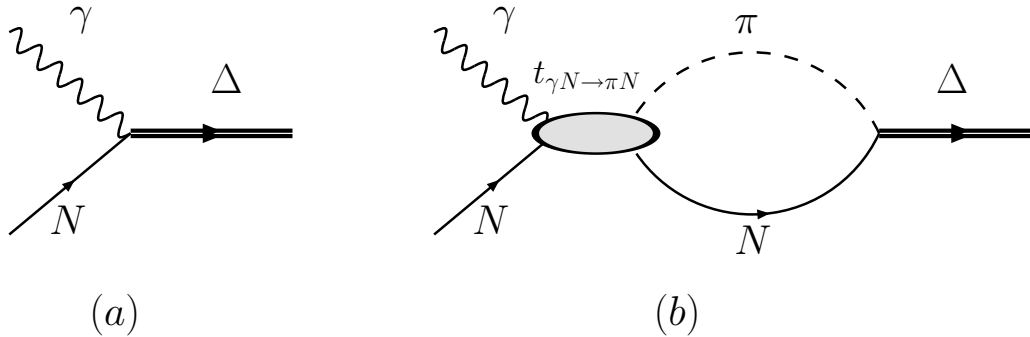
<http://ebac-theory.jlab.org>

in studying πN interactions in a full coupled channel formalism

The main aim is to elucidate the structure of nucleon resonances in **photo** and **electroproduction** processes

- Need to fit several parameters with some constraints
- Use χ^2 minimization
- Our function takes minutes to evaluate
- Thus: we decided to parallelize

Brief on the physics



Relevant points:

- Build a reaction theory to analyze the data
- Including most known phenomenology
- Coupled channel effects

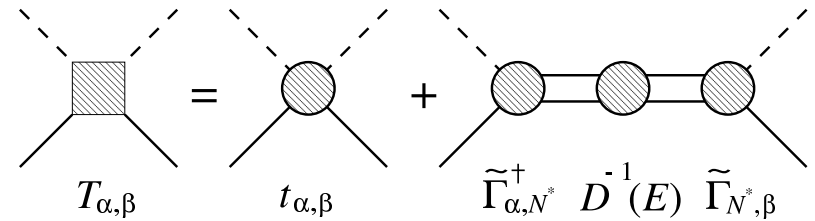
Brief on the physics (II)

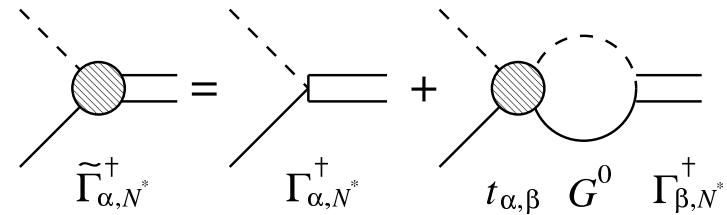
Non-resonant + resonant

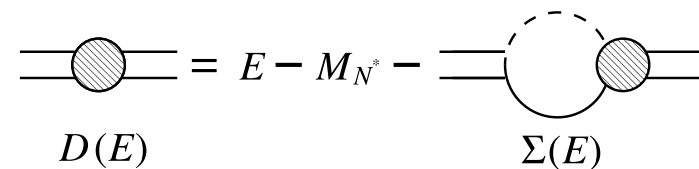
dressed resonant vertexes

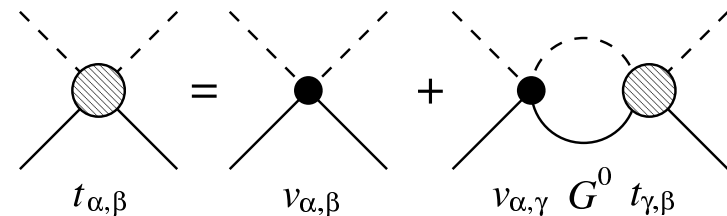
resonant self-energies

non-resonant tmatrix

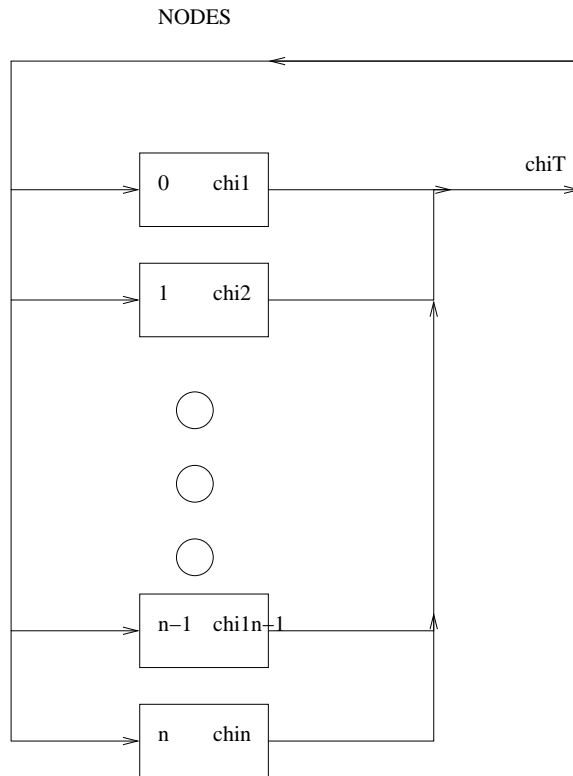
$$T_{\alpha,\beta} = t_{\alpha,\beta} + \tilde{\Gamma}_{\alpha,N^*}^+ \bar{D}^{-1}(E) \tilde{\Gamma}_{N^*,\beta}$$


$$\tilde{\Gamma}_{\alpha,N^*}^+ = \Gamma_{\alpha,N^*}^+ + t_{\alpha,\beta} G^0 \Gamma_{\beta,N^*}^+$$


$$D(E) = E - M_{N^*} - \Sigma(E)$$


$$t_{\alpha,\beta} = v_{\alpha,\beta} + v_{\alpha,\gamma} G^0 t_{\gamma,\beta}$$


Minimizating in parallel



- Our aim was to build a Parallel code
- by minimally changing an existing code

Coupled channel parallel code

The coupled channel code was parallelized: **CCEBA**
and tested at the NERSC supercomputing facility (LBNL)

- At **BSC** we have the project:
“Dynamical Coupled Channel Analysis of Excited
Baryons”
with 100000 hours allotted.

Applying to BSC/CNS

The fifth fastest supercomputer is HERE
and sometimes it even runs...

- <http://www.bsc.es>
- application is EASY and on the WEB (next in April)
- Now each application is for 4 months
- Ask Assum or myself if you need help to fill the application

Some useful tutorials

some good tutorials and really useful guides can be found in
<http://www.nersc.gov>

- <http://www.nersc.gov/nusers/help/tutorials/mpi/intro/>
- http://www.llnl.gov/computing/tutorials/parallel_comp

Summary & Outlook

- We have a nice supercomputer in Barcelona
- and a cluster in the department
- Lets use them! (properly)